1. For which of the following tasks might K-means clustering be a suitable algorithm? Select all that apply. 1 point

- [ ] Given historical weather records, predict if tomorrow's weather will be sunny or rainy.
- [x] From the user usage patterns on a website, figure out what different groups of users exist.
- [x] Given a set of news articles from many different news websites, find out what are the main topics covered.
- [ ] Given many emails, you want to determine if they are Spam or Non-Spam emails.

2. Suppose we have three cluster centroids $\mu_1 = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$, $\mu_2 = \begin{bmatrix} -3 \\ 0 \end{bmatrix}$ and $\mu_3 = \begin{bmatrix} 4 \\ 2 \end{bmatrix}$. Furthermore, we have a training example $x^{(i)} = \begin{bmatrix} -2 \\ 1 \end{bmatrix}$. After a cluster assignment step, what will $c^{(i)}$ be? 1 point

- (•) $c^{(i)} = 2$
- ( ) $c^{(i)} = 3$
- ( ) $c^{(i)}$ is not assigned
- ( ) $c^{(i)} = 1$

3. K-means is an iterative algorithm, and two of the following steps are repeatedly carried out in its inner-loop. Which two? 1 point

- [x] Move the cluster centroids, where the centroids μ_k are updated.
- [ ] Test on the cross-validation set.
- [x] The cluster assignment step, where the parameters $c^{(i)}$ are updated.
- [ ] Randomly initialize the cluster centroids.

4. Suppose you have an unlabeled dataset $\{x^{(1)}, \ldots, x^{(m)}\}$. You run K-means with 50 different random initializations, and obtain 50 different clusterings of the data. What is the recommended way for choosing which one of these 50 clusterings to use? 1 point

- (•) Compute the distortion function $J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_k)$, and pick the one that minimizes this.
- ( ) Use the elbow method.
- ( ) Plot the data and the cluster centroids, and pick the clustering that gives the most "coherent" cluster centroids.
- ( ) Manually examine the clusterings, and pick the best one.

5. Which of the following statements are true? Select all that apply. 1 point

- [x] A good way to initialize K-means is to select K (distinct) examples from the training set and set the cluster centroids equal to these selected examples.
- [ ] K-Means will always give the same results regardless of the initialization of the centroids.
- [ ] Once an example has been assigned to a particular centroid, it will never be reassigned to another different centroid
- [x] On every iteration of K-means, the cost function $J(c^{(1)}, \ldots, c^{(m)}, \mu_1, \ldots, \mu_k)$ (the distortion function) should either stay the same or decrease; in particular, it should not increase.